

# Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see

<https://creativecommons.org/licenses/by-sa/2.0/legalcode>




deeplearning.ai

# Conditional Generation: Intuition

# Outline

- Unconditional generation
- Conditional vs. unconditional generation



# Unconditional Generation

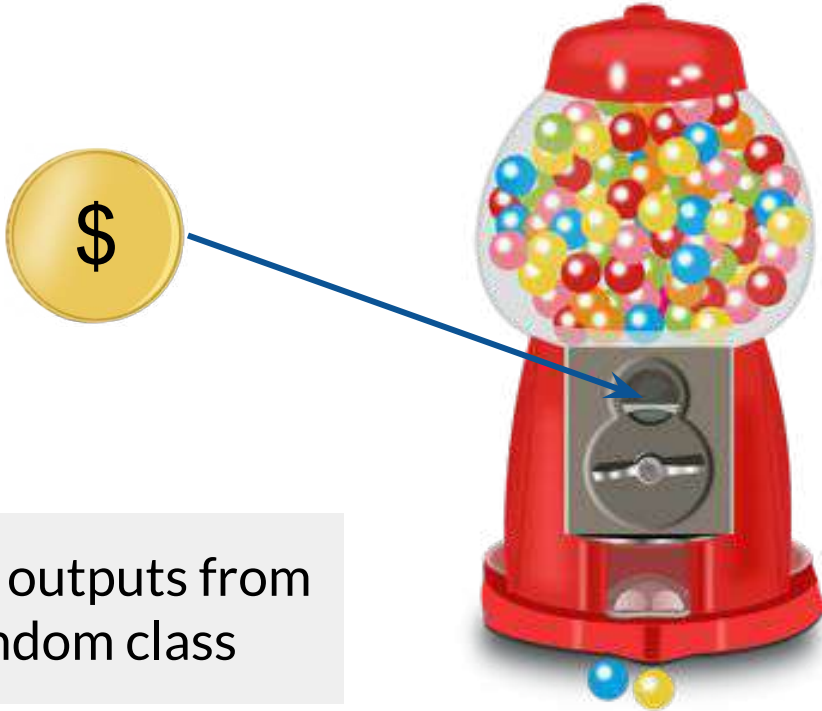
You get outputs from  
a random class

# Unconditional Generation

You get outputs from  
a random class



# Unconditional Generation

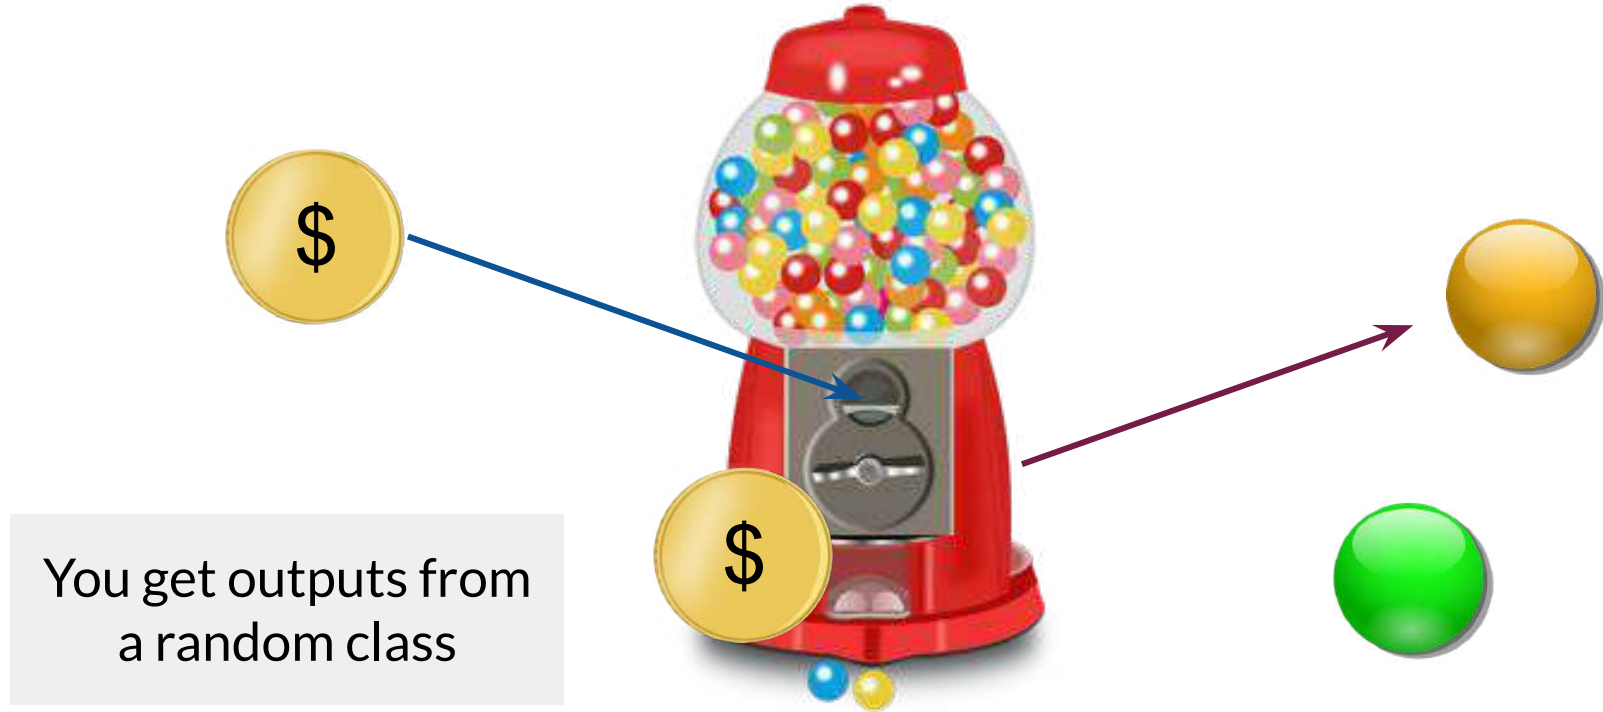


You get outputs from  
a random class

# Unconditional Generation




# Unconditional Generation





# Unconditional Generation




# Unconditional Generation



You get outputs from  
a random class


# Unconditional Generation



# Conditional Generation


You get what you ask for

# Conditional Generation




You get what you ask for

# Conditional Generation



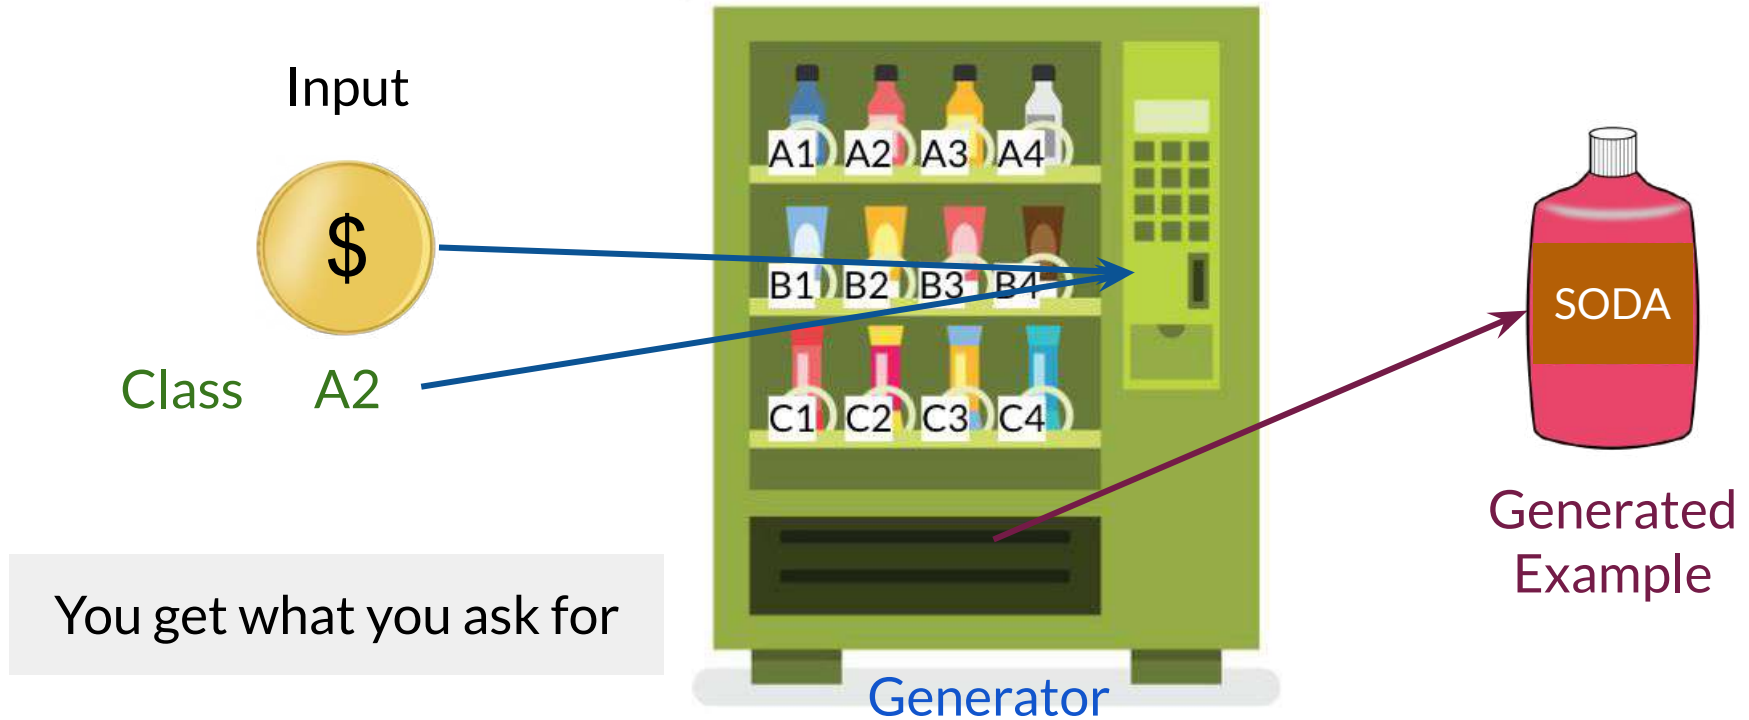
You get what you ask for

# Conditional Generation



You get what you ask for

# Conditional Generation





# Conditional vs. Unconditional Generation



# Conditional vs. Unconditional Generation

Conditional

Unconditional

Examples from **the classes  
you want**


Examples from *random classes*

# Conditional vs. Unconditional Generation

Conditional	Unconditional
Examples from <b>the classes you want</b>	Examples from <i>random classes</i>
Training dataset needs to be <b>labeled</b>	Training dataset <i>doesn't need to be labeled</i>

# Summary

- Conditional generation requires labeled datasets
- Examples can be generated for the selected class






deeplearning.ai


# Conditional Generation: Inputs

# Outline

- How to tell the generator what type of example to produce
- Input representation for the discriminator




# Generator Input




# Generator Input

Noise vector




Class (one-hot) vector






# Generator Input

Noise vector




Class (one-hot) vector



Husky


# Generator Input

Noise vector



Randomness in the  
generation


Class (one-hot) vector



Husky


# Generator Input

Noise vector



Randomness in the generation


Class (one-hot) vector




Husky

Control in the generation

# Generator Input




# Generator Input



# Generator Input

Output




# Discriminator Input




Discriminator

# Discriminator Input






# Discriminator Input



# Discriminator Input



# Discriminator Input



Golden Retriever




Output

98%  
Real


# Discriminator Input




# Discriminator Input



# Discriminator Input



# Discriminator Input



# Summary

- The class is passed to the generator as one-hot vectors
- The class is passed to the discriminator as one-hot matrices
- The size of the vector and the number of matrices represent the number of classes








deeplearning.ai

# Controllable Generation

# Outline


- What is controllable generation
- How it compares to conditional generation



# Controllable Generation

Change specific features of the output


# Controllable Generation



Change specific features of the output

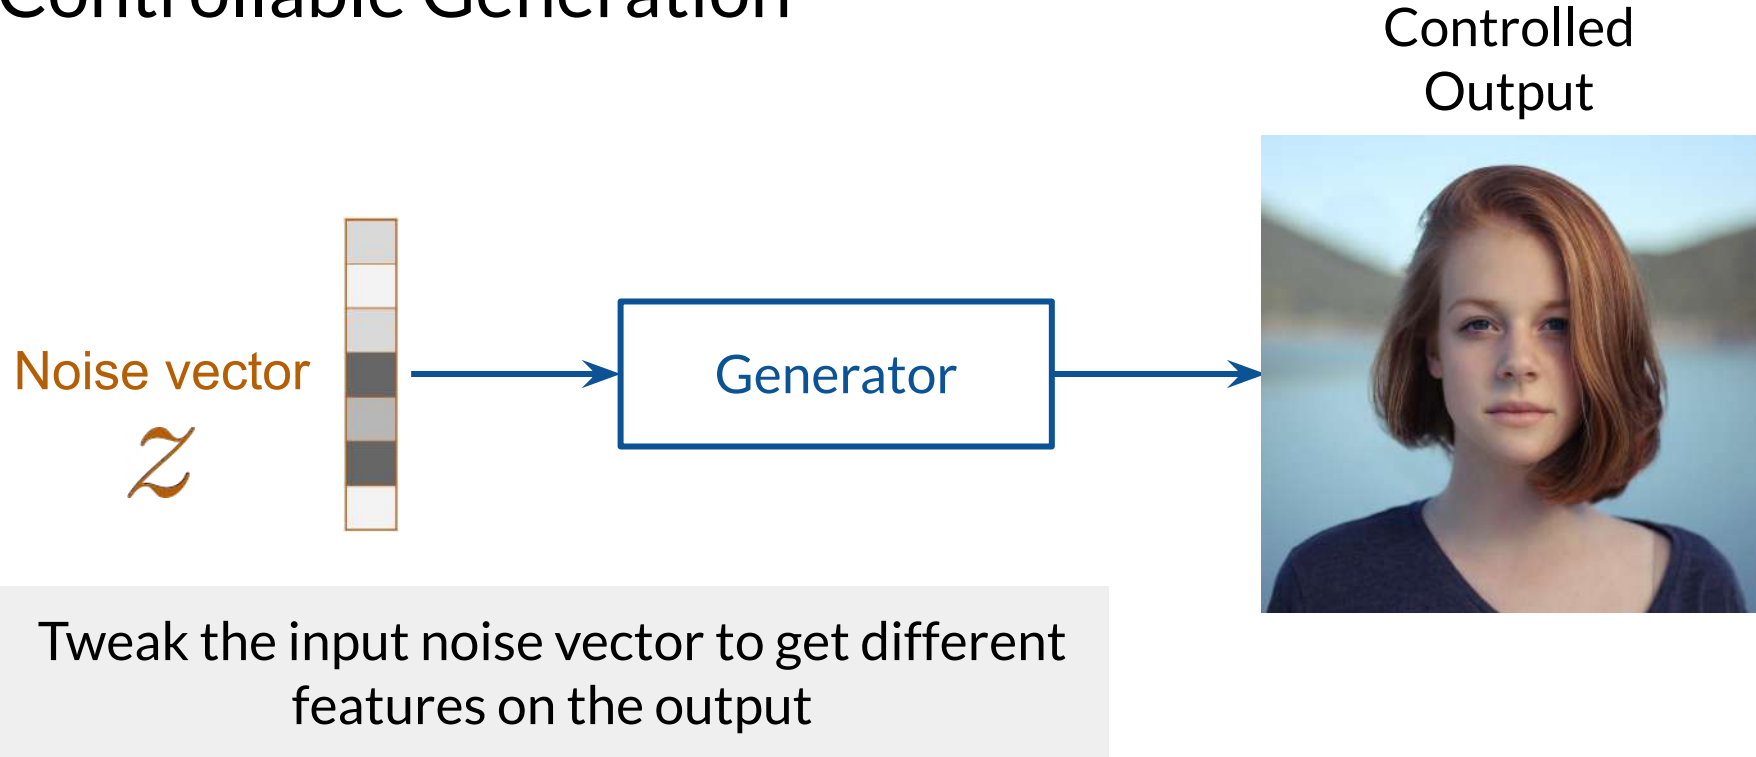
Available from: <https://arxiv.org/abs/1907.10786>

# Controllable Generation




Tweak the input noise vector to get different features on the output

# Controllable Generation



# Controllable Generation



# Controllable Generation vs. Conditional Generation

Controllable

Conditional



# Controllable Generation vs. Conditional Generation

Controllable

Conditional

Examples with the **features**  
**that you want**

Examples from *the classes you*  
*want*

# Controllable Generation vs. Conditional Generation

## Controllable

Examples with the **features  
that you want**

Training dataset **doesn't need  
to be labeled**

## Conditional

Examples from *the classes you  
want*

Training dataset *needs to be  
labeled*

# Controllable Generation vs. Conditional Generation

## Controllable

Examples with the **features that you want**

Training dataset **doesn't need to be labeled**

**Manipulate the  $z$  vector**  
input

## Conditional


Examples from *the classes you want*

Training dataset *needs to be labeled*

*Append a class vector* to the  
input

# Summary

- Controllable generation lets you control the features of the generated outputs
- It does not need a labeled training dataset
- The input vector is tweaked to get different features on the output






deeplearning.ai



# Vector Algebra in the Z-Space

# Outline

- Interpolation in the Z-space
- Modifying the noise vector  $z$  to control desired features




# Interpolation Using the Z-Space




How an image morphs into another

# Interpolation Using the Z-Space





# Interpolation Using the Z-Space




$g(v_1)$



$g(v_2)$



# Interpolation Using the Z-Space




$g(v_1)$



$g(v_2)$



# Interpolation Using the Z-Space




Intermediate images using  
the Z-space


$g(v_1)$   $\longleftrightarrow$   $g(v_2)$



# Z-Space and Controllable Generation






# Z-Space and Controllable Generation




Output

# Z-Space and Controllable Generation




Output

# Z-Space and Controllable Generation




# Z-Space and Controllable Generation





# Z-Space and Controllable Generation




Original output

$$g(v_1) \longrightarrow$$



# Z-Space and Controllable Generation




Original output

$$g(v_1) \longrightarrow$$



# Z-Space and Controllable Generation



Original output

$$g(v_1) \rightarrow$$




Controlled output

$$g(v_1 + d) \rightarrow$$



# Summary

- To control output features, you need to find directions in the Z-space
- To modify your output, you move around in the Z-space





deeplearning.ai

# Challenges with Controllable Generation

# Outline

- Output feature correlation
- Z-space entanglement



# Feature Correlation

Uncorrelated  
Features



Add beard



# Feature Correlation

Uncorrelated  
Features



Add beard




Correlated  
Features



Add beard

Make more  
masculine





# Z-Space Entanglement

Noise vector

$z$



Glasses

Beard


Hair

Eyes


Age

Output  
Features


# Z-Space Entanglement




# Z-Space Entanglement




# Z-Space Entanglement



# Z-Space Entanglement




# Z-Space Entanglement



It is not possible to control single output features

# Z-Space Entanglement



It is not possible to control single output features

# Summary

- When trying to control one feature, others that are correlated change
- Z-space entanglement makes controllability difficult, if not impossible
- Entanglement happens when  $z$  does not have enough dimensions








deeplearning.ai

# Classifier Gradients

# Outline

- How to use classifiers to find directions in the Z-space
- Requirements to use this method




# Classifier Gradients

# Classifier Gradients

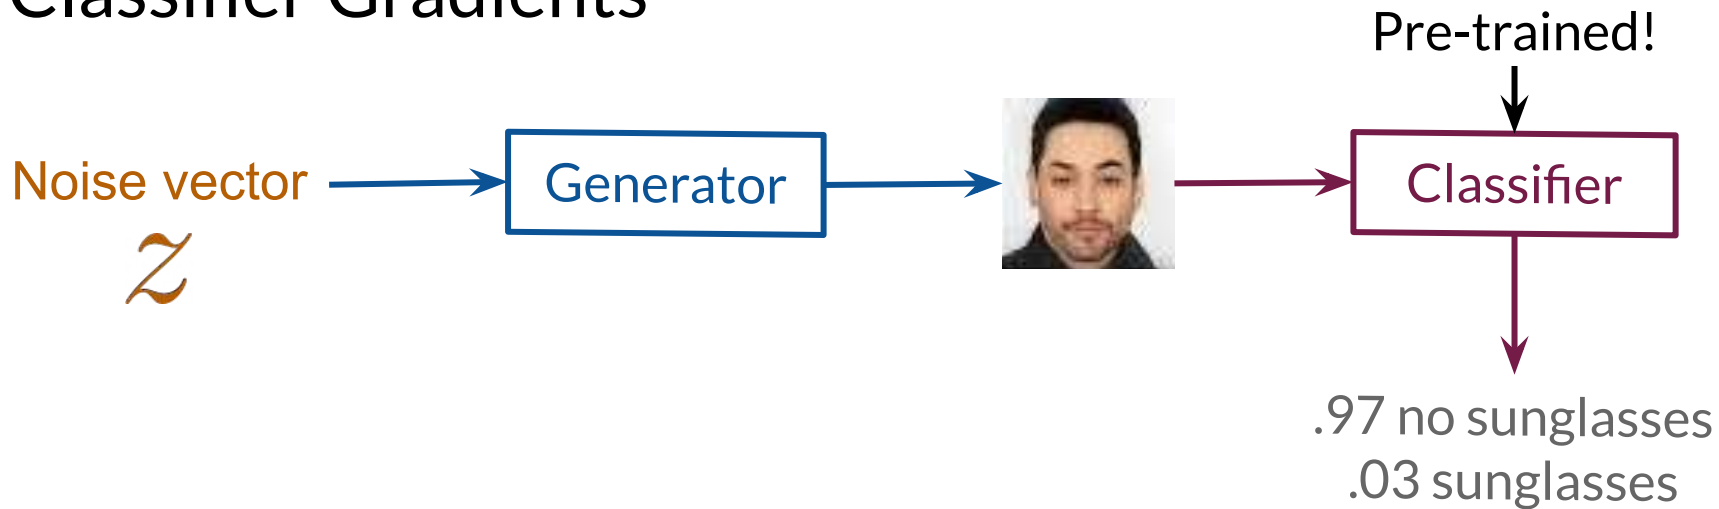
Noise vector

$z$


# Classifier Gradients




# Classifier Gradients



# Classifier Gradients




# Classifier Gradients



Modify **just** the **noise vector** until the feature emerges



# Classifier Gradients



Modify **just** the **noise vector** until the feature emerges

# Summary

- Classifiers can be used to find directions in the Z-space
- To find directions, the updates are done just to the noise vector






deeplearning.ai

# Disentanglement

# Outline


- What a disentangled Z-space means
- Ways to encourage disentangled Z-spaces



# Disentangled Z-Space

$$v_1 = [ 1, 2, 3, \dots ]$$

$$v_2 = [ 5, 6, 7, \dots ]$$




# Disentangled Z-Space

$$v_1 = [ \overset{z_1}{1}, 2, 3, \dots ]$$


$$v_2 = [ 5, 6, 7, \dots ]$$

Hair  
color




# Disentangled Z-Space

$$\begin{array}{c} z_1 \quad z_2 \\ v_1 = [ \text{1}, \text{2}, 3, \dots ] \\ v_2 = [ \text{5}, \text{6}, 7, \dots ] \\ \text{Hair color} \quad \text{Hair length} \end{array}$$



# Disentangled Z-Space

$$\begin{array}{c} z_1 \quad z_2 \\ v_1 = [ \text{1}, \text{2}, 3, \dots ] \\ v_2 = [ \text{5}, \text{6}, 7, \dots ] \\ \text{Hair color} \quad \text{Hair length} \end{array}$$






# Disentangled Z-Space

$$\begin{array}{c} z_1 \quad z_2 \\ v_1 = [ \text{1}, \text{2}, 3, \dots ] \\ v_2 = [ \text{5}, \text{6}, 7, \dots ] \end{array}$$

Hair color      Hair length




# Disentangled Z-Space


$$\begin{array}{c} z_1 \quad z_2 \\ v_1 = [ \text{1}, \text{2}, 3, \dots ] \\ v_2 = [ \text{5}, \text{6}, 7, \dots ] \end{array}$$

Hair color      Hair length


Latent factors of variation



# Disentangled Z-Space




# Disentangled Z-Space



Changes to one feature  
don't affect the others

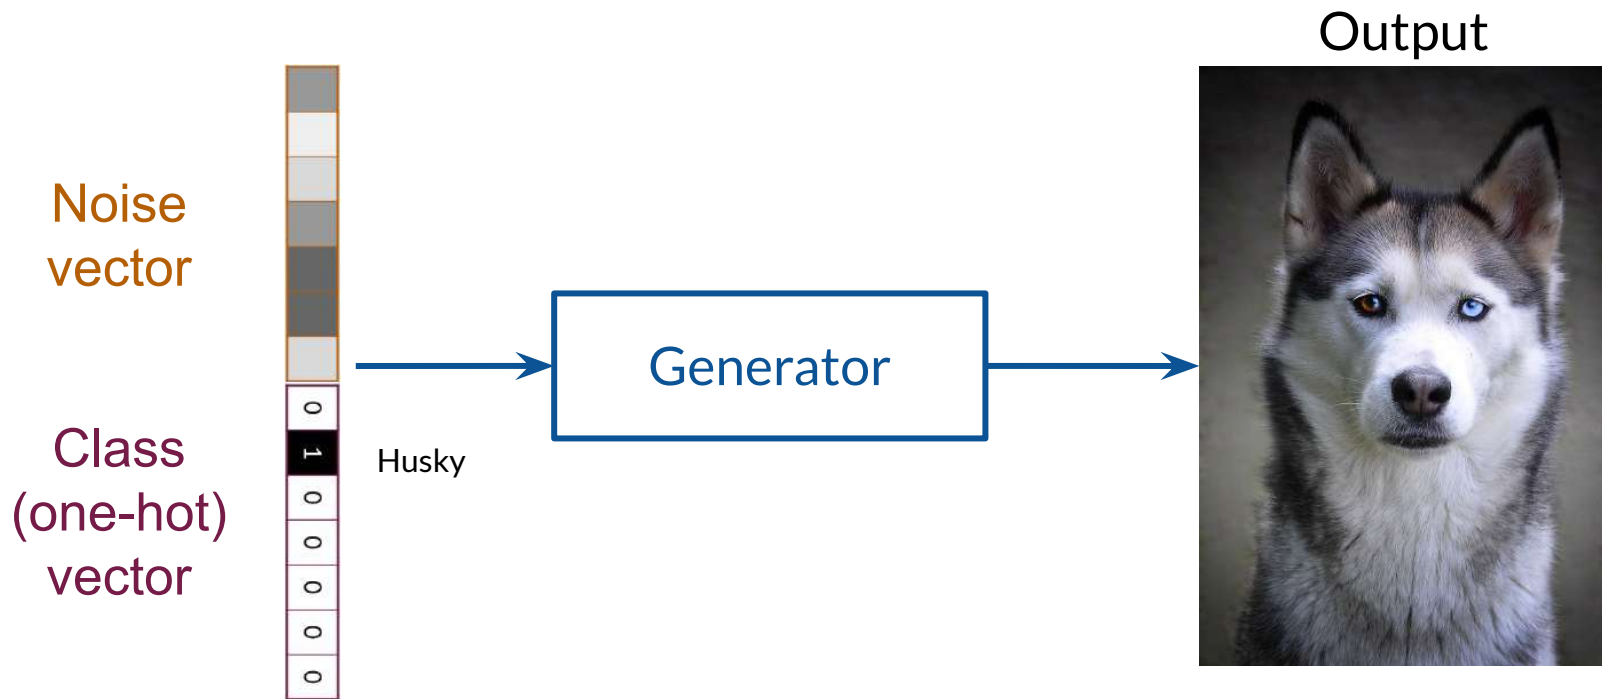
# Disentangled Z-Space



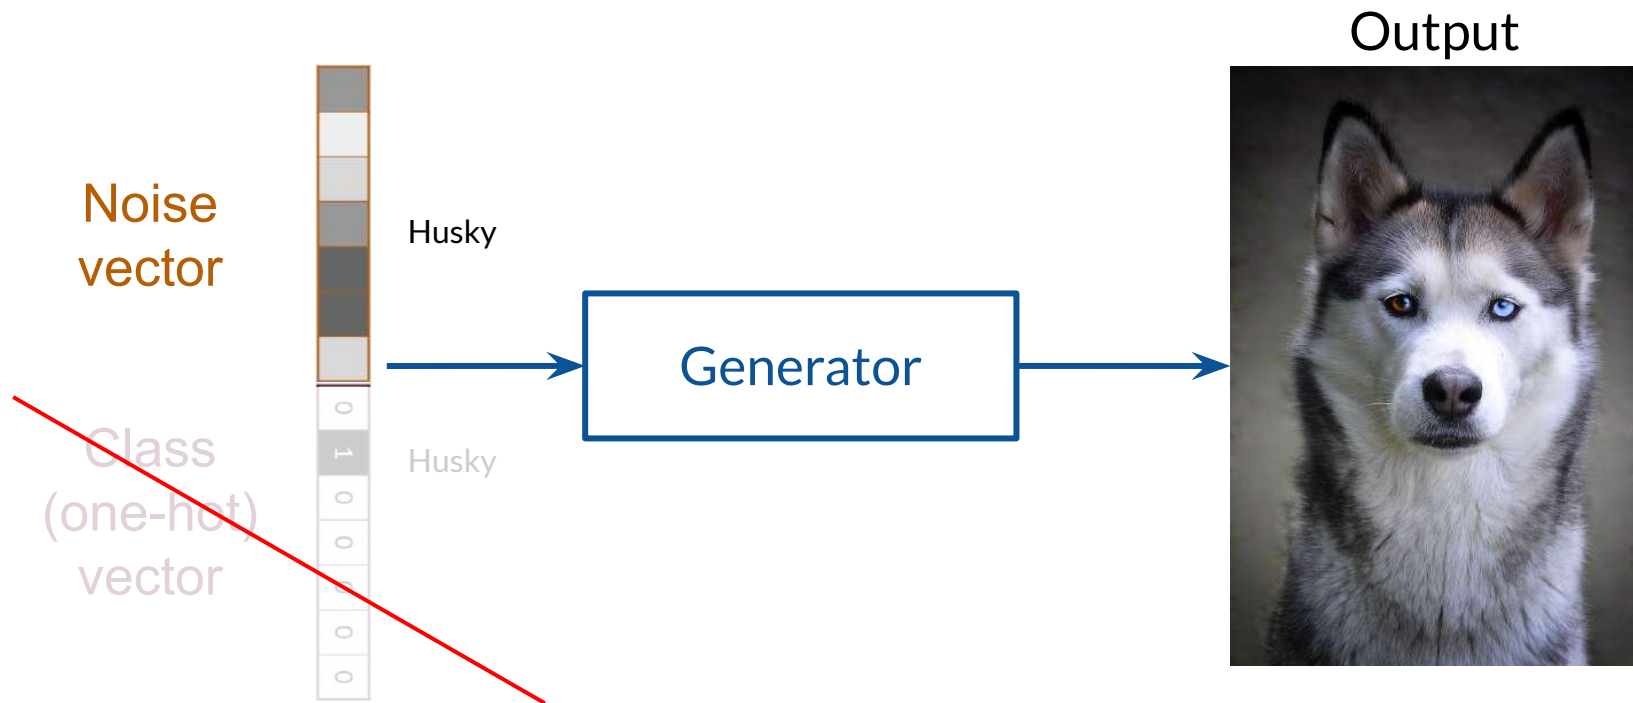
Changes to one feature  
don't affect the others



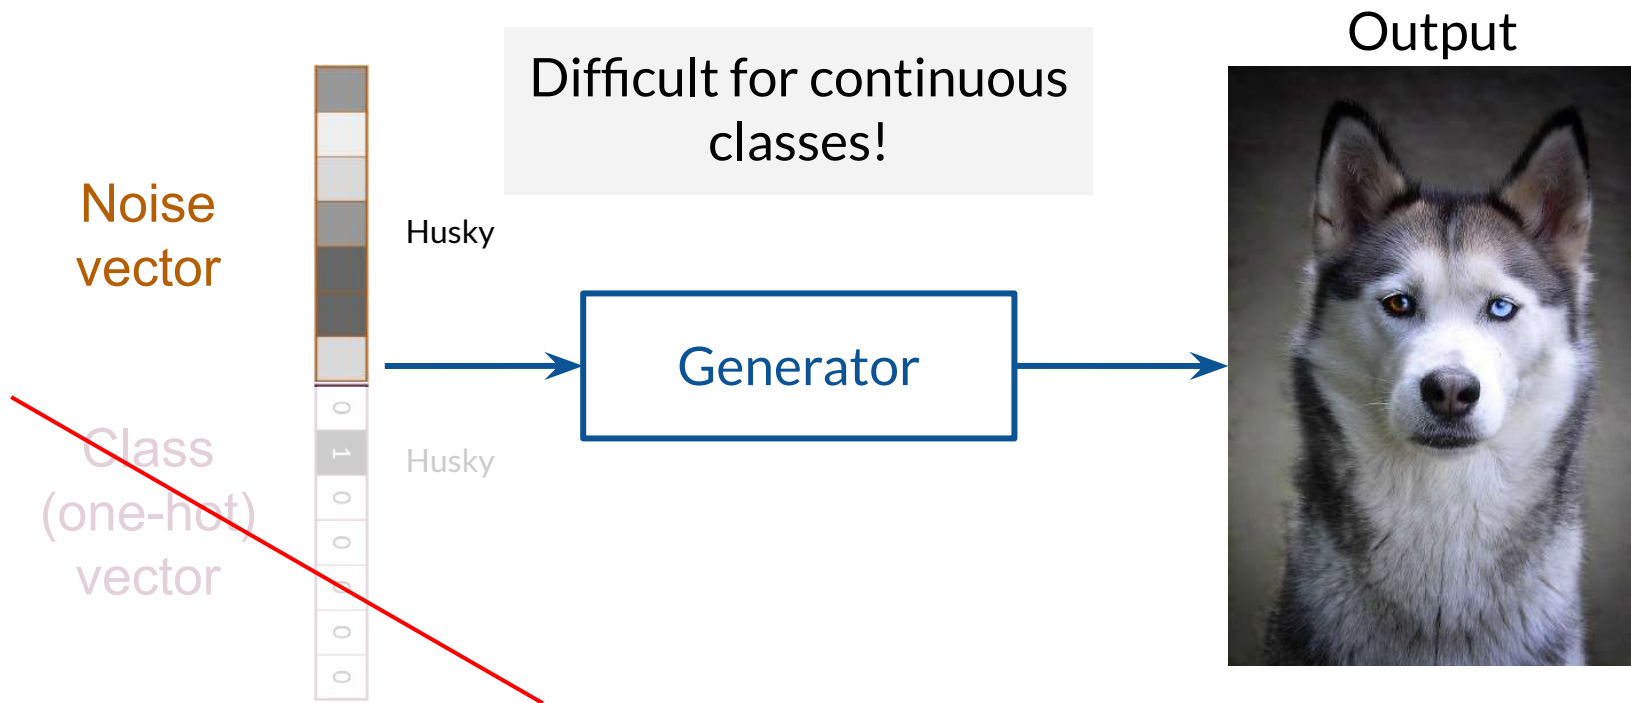
# Encourage Disentanglement: Supervision



# Encourage Disentanglement: Supervision



# Encourage Disentanglement: Supervision





# Encourage Disentanglement: Loss Function

$$v_1 = [ 1, 2, 3, \dots ]$$

$$v_2 = [ 5, 6, 7, \dots ]$$

# Encourage Disentanglement: Loss Function

$$v_1 = [ 1, 2, 3, \dots ]$$

$$v_2 = [ 5, 6, 7, \dots ]$$

$$L_{\text{new}} = \boxed{L_{\text{original}}} + \boxed{\text{reg}_d}$$

Original loss      Regularization

# Encourage Disentanglement: Loss Function

$$v_1 = [1, 2, 3, \dots]$$

$$v_2 = [5, 6, 7, \dots]$$

$$L_{\text{new}} = \boxed{L_{\text{original}}} + \boxed{\text{reg}_d}$$

Original loss      Regularization

Can be any loss function  
(e.g. BCE, W-Loss)

# Encourage Disentanglement: Loss Function


$$\begin{array}{cc} & z_1 & z_2 \\ v_1 = [ & 1, & 2, & 3, \dots ] \\ v_2 = [ & 5, & 6, & 7, \dots ] \end{array}$$

$$L_{\text{new}} = \boxed{L_{\text{original}}} + \boxed{\text{reg}_d}$$

Original loss      Regularization

Can be any loss function  
(e.g. BCE, W-Loss)

# Encourage Disentanglement: Loss Function



$$L_{\text{new}} = \boxed{L_{\text{original}}} + \boxed{\text{reg}_d}$$

Original loss      Regularization

Can be any loss function  
(e.g. BCE, W-Loss)

# Summary

- Disentangled Z-spaces let you control individual features by corresponding  $z$  values directly to them
- There are supervised and unsupervised methods to achieve disentanglement

